

Paddy Prediction

The training set has images **and the following metadata**

Loaded CSV: data/train.csv (rows=10407, cols=4)


	image_id	label	variety	age
0	100330.jpg	bacterial_leaf_blight	ADT45	45
1	100365.jpg	bacterial_leaf_blight	ADT45	45
2	100382.jpg	bacterial_leaf_blight	ADT45	45
3	100632.jpg	bacterial_leaf_blight	ADT45	45
4	101918.jpg	bacterial_leaf_blight	ADT45	45

The submission set (in /test_images), has **just an image, no metadata**

Using the metadata during training overview

Cannot feed the metadata to the model during training because its not there when inferenceing on submission set.

But... Models often perform better on everything when number of tasks increases. So get the model to predict the metadata during training **with only the image as input to the model (so can still inference with just image)**. Instead of 1 head, add 3. W_1, W_2 and W_3 weight each loss.



Need

- Dataset that can deliver the image and label, variety and age metadata
- model with 3 heads
- criterion (loss function) that can do loss on each head and produce a weighted sum
- plus various changes to code (llms will automatically do this)

Strategy:

train the FC layers first with the rest of the model frozen
then unfreeze entire model and train some more at $lr/10$ **(this is where model will ingest new information that comes from it predicting age and variety)**